

You can access this page also inside the Remote Desktop by using the icons on the desktop

- Score
- Questions and Answers
- Preview Questions and Answers
- Exam Tips

CKAD Simulator Kubernetes 1.33

<https://killer.sh>

Each question needs to be solved on a specific instance other than your main `candidate@terminal`. You'll need to connect to the correct instance via ssh, the command is provided before each question. To connect to a different instance you always need to return first to your main terminal by running the `exit` command, from there you can connect to a different one.

In the real exam each question will be solved on a different instance whereas in the simulator multiple questions will be solved on same instances.

Question 1 | Namespaces

Solve this question on instance: `ssh ckad5601`

The DevOps team would like to get the list of all *Namespaces* in the cluster.

The list can contain other columns like STATUS or AGE.

Save the list to `/opt/course/1/namespaces` on `ckad5601`.

Answer:

```
k get ns > /opt/course/1/namespaces
```

The content should then look like:

```
# /opt/course/1/namespaces
```

NAME	STATUS	AGE
default	Active	136m
earth	Active	105m
jupiter	Active	105m
kube-node-lease	Active	136m
kube-public	Active	136m
kube-system	Active	136m
mars	Active	105m
shell-intern	Active	105m

Question 2 | Pods

Solve this question on instance: `ssh ckad5601`

Create a single *Pod* of image `httpd:2.4.41-alpine` in *Namespace* `default`. The *Pod* should be named `pod1` and the container should be named `pod1-container`.

Your manager would like to run a command manually on occasion to output the status of that exact *Pod*. Please write a command that does this into `/opt/course/2/pod1-status-command.sh` on `ckad5601`. The command should use `kubect1`.

Answer:

```
k run # help
```

```
k run pod1 --image=httpd:2.4.41-alpine --dry-run=client -oyaml > 2.yaml
```

```
vim 2.yaml
```

Change the container name in `2.yaml` to `pod1-container`:

```
# 2.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: pod1
  name: pod1
spec:
  containers:
    - image: httpd:2.4.41-alpine
      name: pod1-container # change
      resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

Then run:

```
→ k create -f 2.yaml
```

```
pod/pod1 created
```

```
→ k get pod
```

NAME	READY	STATUS	RESTARTS	AGE
pod1	0/1	ContainerCreating	0	6s

```
→ k get pod
```

NAME	READY	STATUS	RESTARTS	AGE
pod1	1/1	Running	0	30s

Next create the requested command:

```
vim /opt/course/2/pod1-status-command.sh
```

The content of the command file could look like:

```
# /opt/course/2/pod1-status-command.sh
kubectl -n default describe pod pod1 | grep -i status:
```

Another solution would be using jsonpath:

```
# /opt/course/2/pod1-status-command.sh
kubectl -n default get pod pod1 -o jsonpath="{.status.phase}"
```

To test the command:

```
→ sh /opt/course/2/pod1-status-command.sh
Running
```

Question 3 | Job

Solve this question on instance: `ssh ckad7326`

Team Neptune needs a *Job* template located at `/opt/course/3/job.yaml`. This *Job* should run image `busybox:1.31.0` and execute `sleep 2 && echo done`. It should be in namespace `neptune`, run a total of 3 times and should execute 2 runs in parallel.

Start the *Job* and check its history. Each pod created by the *Job* should have the label `id: awesome-job`. The job should be named `neb-new-job` and the container `neb-new-job-container`.

Answer:

```
k -n neptune create job -h
```

```
k -n neptune create job neb-new-job --image=busybox:1.31.0 --dry-run=client -oyaml -- sh -c "sleep 2 && echo done" > /opt/course/3/job.yaml
```

```
vim /opt/course/3/job.yaml
```

Make the required changes in the yaml:

```
# /opt/course/3/job.yaml
apiVersion: batch/v1
kind: Job
metadata:
  creationTimestamp: null
  name: neb-new-job
  namespace: neptune
spec:
  completions: 3          # add
  parallelism: 2          # add
  template:
    metadata:
      creationTimestamp: null
      labels:              # add
        id: awesome-job   # add
    spec:
      containers:
      - command:
        - sh
        - -c
        - sleep 2 && echo done
        image: busybox:1.31.0
        name: neb-new-job-container # update
        resources: {}
      restartPolicy: Never
status: {}
```

Then to create it:

```
k -f /opt/course/3/job.yaml create # namespace already set in yaml hence not needed
```

Check *Job* and *Pods*, you should see two running parallel at most but three in total:

```
→ k -n neptune get pod,job | grep neb-new-job
```

pod/neb-new-job-jhq2g	0/1	ContainerCreating	0	4s
pod/neb-new-job-vf6ts	0/1	ContainerCreating	0	4s

job.batch/neb-new-job	0/3	4s	5s
-----------------------	-----	----	----

```
→ k -n neptune get pod,job | grep neb-new-job
```

pod/neb-new-job-gm8sz	0/1	ContainerCreating	0	0s
pod/neb-new-job-jhq2g	0/1	Completed	0	10s
pod/neb-new-job-vf6ts	1/1	Running	0	10s

job.batch/neb-new-job	1/3	10s	11s
-----------------------	-----	-----	-----

```
→ k -n neptune get pod,job | grep neb-new-job
```

pod/neb-new-job-gm8sz	0/1	ContainerCreating	0	5s
-----------------------	-----	-------------------	---	----

pod/neb-new-job-jhq2g	0/1	Completed	0	15s
pod/neb-new-job-vf6ts	0/1	Completed	0	15s

job.batch/neb-new-job	2/3	15s	16s
-----------------------	-----	-----	-----

```
→ k -n neptune get pod,job | grep neb-new-job
```

pod/neb-new-job-gm8sz	0/1	Completed	0	12s
pod/neb-new-job-jhq2g	0/1	Completed	0	22s

pod/neb-new-job-vf6ts	0/1	Completed	0	22s
job.batch/neb-new-job	3/3	21s	23s	

Check history:

```
→ k -n neptune describe job neb-new-job
...
```

Events:				
Type	Reason	Age	From	Message
----	-----	----	----	-----
Normal	SuccessfulCreate	2m52s	job-controller	Created pod: neb-new-job-jhq2g
Normal	SuccessfulCreate	2m52s	job-controller	Created pod: neb-new-job-vf6ts
Normal	SuccessfulCreate	2m42s	job-controller	Created pod: neb-new-job-gm8sz

At the age column we can see that two `Pods` run parallel and the third one after that. Just as it was required in the task.

Question 4 | Helm Management

Solve this question on instance: `ssh ckad7326`

Team Mercury asked you to perform some operations using Helm, all in *Namespace* `mercury`:

1. Delete release `internal-issue-report-apiv1`
2. Upgrade release `internal-issue-report-apiv2` to any newer version of chart `killershell/nginx` available
3. Install a new release `internal-issue-report-apache` of chart `killershell/apache`. The *Deployment* should have two replicas, set these via Helm-values during install
4. There seems to be a broken release, stuck in `pending-install` state. Find it and delete it

Answer:

Helm Chart: Kubernetes YAML template-files combined into a single package, *Values* allow customisation

Helm Release: Installed instance of a *Chart*

Helm Values: Allow to customise the YAML template-files in a *Chart* when creating a *Release*

Step 1

First we should delete the required release:

```
→ helm -n mercury ls
```

NAME	NAMESPACE	...	STATUS	CHART
internal-issue-report-apiv1	mercury	...	deployed	nginx-18.1.14
internal-issue-report-apiv2	mercury	...	deployed	nginx-18.1.14
internal-issue-report-app	mercury	...	deployed	nginx-18.1.14

```
→ helm -n mercury uninstall internal-issue-report-apiv1
release "internal-issue-report-apiv1" uninstalled

→ helm -n mercury ls
```

NAME	NAMESPACE	...	STATUS	CHART
internal-issue-report-apiv2	mercury	...	deployed	nginx-18.1.14
internal-issue-report-app	mercury	...	deployed	nginx-18.1.14

Step 2

Next we need to upgrade a release, for this we could first list the charts of the repo:

```
→ helm repo list
```

NAME	URL
killershell	http://localhost:6000

```
→ helm repo update
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "killershell" chart repository
Update Complete. *Happy Helming!*

→ helm search repo nginx --versions
```

NAME	CHART VERSION	DESCRIPTION
killershell/nginx	18.2.0	NGINX Open Source is a...
killershell/nginx	18.1.15	NGINX Open Source is a...
killershell/nginx	18.1.14	NGINX Open Source is a...
killershell/nginx	13.0.0	NGINX Open Source is a...

Here we see that two newer chart versions are available. But the question only requires us to upgrade to any newer chart version available, so we can simply run:

```
→ helm -n mercury upgrade internal-issue-report-apiv2 killershell/nginx
Release "internal-issue-report-apiv2" has been upgraded. Happy Helming!
NAME: internal-issue-report-apiv2
LAST DEPLOYED: Mon Aug 25 14:21:24 2025
NAMESPACE: mercury
STATUS: deployed
REVISION: 2
TEST SUITE: None

→ helm -n mercury ls
```

NAME	NAMESPACE	...	STATUS	CHART	APP VERSION
internal-issue-report-apiv2	mercury	...	deployed	nginx-18.2.0	1.27.1
internal-issue-report-app	mercury	...	deployed	nginx-18.1.14	1.27.1

Looking good!

INFO: Also check out `helm rollback` for undoing a helm rollout/upgrade

Step 3

Now we're asked to install a new release, with a customised values setting. For this we first list all possible value settings for the chart, we can do this via:

```
→ helm show values killershell/apache
global:
  imageRegistry: ""
  imagePullSecrets: []
kubeVersion: ""
nameOverride: ""
fullnameOverride: ""
commonLabels: {}
commonAnnotations: {}
extraDeploy: []
image:
  registry: docker.io
  repository: httpd
  pullPolicy: IfNotPresent
  pullSecrets: []
  debug: false
replicaCount: 1
revisionHistoryLimit: 10
podAffinityPreset: ""
podAntiAffinityPreset: soft
extraPodSpec: {}
```

Or to parse yaml and render with colors:

```
helm show values killershell/apache | yq e
```

This can be a huge list for larger Helm charts. We should find the setting `replicaCount: 1` on top level. This means we can run:

```
→ helm -n mercury install internal-issue-report-apache killershell/apache --set replicaCount=2
NAME: internal-issue-report-apache
LAST DEPLOYED: Mon Aug 25 14:23:38 2025
NAMESPACE: mercury
STATUS: deployed
REVISION: 1
TEST SUITE: None
```

If we would also need to set a value on a deeper level, for example `image.debug`, we could run:

```
helm -n mercury install internal-issue-report-apache killershell/apache \
--set replicaCount=2 \
--set image.debug=true
```

Install done, let's verify what we did:

```
→ helm -n mercury ls
```

NAME	NAMESPACE	...	STATUS	CHART
internal-issue-report-apache	mercury	...	deployed	apache-11.2.20
internal-issue-report-apiv2	mercury	...	deployed	nginx-18.2.0
internal-issue-report-app	mercury	...	deployed	nginx-18.1.14

```
→ k -n mercury get deploy internal-issue-report-apache
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
internal-issue-report-apache	2/2	2	2	64s

We see a healthy deployment with two replicas!

Step 4

By default releases in `pending-upgrade` state aren't listed, but we can show all to find and delete the broken release:

```
→ helm -n mercury ls -a
```

NAME	NAMESPACE	...	STATUS	CHART
internal-issue-report-apache	mercury	...	deployed	apache-11.2.20
internal-issue-report-apiv2	mercury	...	deployed	nginx-18.2.0
internal-issue-report-app	mercury	...	deployed	nginx-18.1.14
internal-issue-report-daniel	mercury	...	pending-install	nginx-18.1.14

```
→ helm -n mercury uninstall internal-issue-report-daniel  
release "internal-issue-report-daniel" uninstalled
```

Thank you Helm for making our lives easier! (Till something breaks)

Question 5 | ServiceAccount, Secret

Solve this question on instance: `ssh ckad7326`

Team Neptune has its own *ServiceAccount* named `neptune-sa-v2` in *Namespace* `neptune`. A coworker needs the token from the *Secret* that belongs to that *ServiceAccount*. Write the base64 decoded token to file `/opt/course/5/token` on `ckad7326`.

Answer:

Secrets won't be created automatically for **ServiceAccounts*, but it's possible to create a *Secret* manually and attach it to a *ServiceAccount* by setting the correct annotation on the *Secret*. This was done for this task.

```
k -n neptune get sa # get overview  
k -n neptune get secrets # shows all secrets of namespace  
k -n neptune get secrets -oyaml | grep annotations -A 1 # shows secrets with first annotation
```

If a *Secret* belongs to a *ServiceAccount*, it'll have the annotation `kubernetes.io/service-account.name`. Here the *Secret* we're looking for is `neptune-secret-1`.


```
→ k -n neptune get secret neptune-secret-1 -o yaml
apiVersion: v1
data:
...
  token:
ZXlKaGJHY2lPaUpTVXpJMU5pSXNjbXRwWkNjNk1tNWFaRmRxWkRKMMFHTnZRM0JxV0haTlIxZzZfiM3BJY201S1owaEhOV3hUWmt3elFuRmFhV
EZhZDJNaWZRLmV5SnBjM01pT2lKcmRXSmxjbTVsZEdWekwzTmxjblpwWTJWafkyTnZkVzUwSWl3aWEzVmlaWEplWlhSbGN5NXBieTl6WlhKMM
FXImxZV05qYjNWdWRDOXVZVzFsYzNCAfkyVWlPaUp1WlhCMGRXNWxJaXdpYTnWVaWpYSnVaWFJsY3k1cGJ5OXpaWEoyYVdObFlXTmpiMlZlZEM
5elpXTnlaWFFlYm1GdFpTSTZJbTVsY0hSMWJtVXRjMkV0ZGpJdGRHOXJaVzR0Wm5FNU1tb2lMQ0pyZFdKbGNtNWxkR1Z6TG1sdkwzTmxjblpw
WTJWafkyTnZkVzUwTDNOBGnuWnBZMlV0WVdOamIzVnVkJzV1WVcxblqb2libVZ3ZEhWdVpTMXpZUzEyTWlJc0ltdDFZbVZ5Ym1WMMFpYTXVhV
zh2YzJWeWRtbGpaV0ZqWTI5MWJuUXZjMlZ5ZG1salpTMWhZMk52ZFclMExuVnBaQ0k2SWpZM1ltUmpOak0yTFRKbFl6TXROREpoWkMwNE9HRT
FMV0ZoWxpGbfqWmxPVFpsTlNjC0luTjFZaUk2SW5ONWMzUmxiVHB6WlhKMmFXTmxZV05qYjNWdWREcHVaWEIwZFclbE9tNWxjSF1xYm1VdGM
yRXRkaklpZlEuV1lnYm9NNENUZDBwZENKNzh3alV3bXRhbGgtMnZzS2pBTnlQc2gtNmd1RXdpdFdcTVGYnc1WkhQdHZBZBJmBFB6cE9IRWJB
ZTRlVU05NUJSR1diWUlkZDp1TjklSjBENFJORmtWVXQ0OHR3b2FrUlY3aC1hUHV3c1FYSGhaWnp5NHlplbUZIRz1VZm1zazVZcjRSVmNHNm4xM
zd5LUZIMDhLOHpaaklQQXNLRHFOQlF0eGctbFp2d1ZNaT2aUlocnJ6QVFzMElCT1Y4Mk9KWUd5Mm8tVlFWYzBVVWFuQ2Y5NFkzZlZlQwWVRpcV
F2Y3pZTXM2bno5dXQtWGD3aXRyQlk2VG05QmdQcHJBOWtfaJvXRXhfTFVVVWlVwUEFpRU43T3pka0pzSThjDHRoMTBseXBjMUFlRnI0M3Q2QUx
5clFvQk0zOWFiRGZxM0Zrc1Itb2NfV013
kind: Secret
...
```

This shows the base64 encoded token. To get the decoded one we could pipe it manually through `base64 -d` or we simply do:

```
→ k -n neptune describe secret neptune-secret-1
...
Data
=====
token:
eyJhbGciOiJSUzI1NiIsImtpZCI6Im5aZFdqZDJ2aGNvQ3BqWHZORlglb3pIcm5JZ0hHNWxTZkxzQnFaaTFad2MifQ.eyJpc3MiOiJrdWJlcm
5ldGVzL3NlcnZpY2VhY2NvdW50Iiwia3ViZXJuZXRlcy5pby9zZXJ2aWNlYWNjb3VudC9uYW1lc3BhY2UiOiJuZXB0dW5lIiwia3ViZXJuZXR
lcy5pby9zZXJ2aWNlYWNjb3VudC9zZWNYZXQubmFtZSI6Im5lcHRlbnUtc2EtZjItZDg9rZW4tZnE5Mm0iLCJrdWJlcm5ldGVzLmlvL3NlcnZp
Y2VhY2NvdW50L3NlcnZpY2UtYWNjb3VudC5uYW1lIjoibmVudHVuZS1zYS12MiIsImt1YmVybmV0ZXMuaW8vc2VydmljZWZjY291bnQvc2Vyd
mljZS1hY2NvdW50LnVpZCI6IjY2YmRjNjM2LTJlYzMtNDJhZC00OGE1LWFhYzFlZjZlOTZlNSIsInN1YiI6InN5c3RlbTpzZXJ2aWNlYWNjb3
VudDpuZXB0dW5lOm5lcHRlbnUtc2EtZjIifQ.VYgboM4CTd0pdCJ78wjUwmtalh-2vsKjANyPsh-6guEwOtWEq5Fbw5ZHPtvAdrLlPzpOHEbA
e4eUM95BRGWbYIdwjuN95J0D4RNFkVUt48twoakRV7h-aPuwsQXhHZZzy4yimFHG9Ufmsk5Yr4RVcG6n137y-FH08K8zZjIPAsKDqNBQtXg-
lZvwVMi6viIhrrzAQs0MBOV820JYgy2o-WQVc0UUanCf94Y3gT0YTiQVvczYMs6nz9ut-
XgwitrBY6Tj9BgPprA9k_j5qEx_LUUZUpPAiEN7OzdkJsI8ctth101ypI1AeFr43t6ALyrQoBM39abDfq3FksR-oc_WMw
ca.crt:      1066 bytes
namespace:   7 bytes
```

Copy the token (part under `token:`) and paste it using vim.

```
vim /opt/course/5/token
```

File `/opt/course/5/token` should contain the token:

```
# /opt/course/5/token
eyJhbGciOiJSUzI1NiIsImtpZCI6Im5aZFdqZDJ2aGNvQ3BqWHZORlglb3pIcm5JZ0hHNWxTZkxzQnFaaTFad2MifQ.eyJpc3MiOiJrdWJlcm
5ldGVzL3NlcnZpY2VhY2NvdW50Iiwia3ViZXJuZXRlcy5pby9zZXJ2aWNlYWNjb3VudC9uYW1lc3BhY2UiOiJuZXB0dW5lIiwia3ViZXJuZXR
lcy5pby9zZXJ2aWNlYWNjb3VudC9zZWNYZXQubmFtZSI6Im5lcHRlbnUtc2EtZjItZDg9rZW4tZnE5Mm0iLCJrdWJlcm5ldGVzLmlvL3NlcnZp
Y2VhY2NvdW50L3NlcnZpY2UtYWNjb3VudC5uYW1lIjoibmVudHVuZS1zYS12MiIsImt1YmVybmV0ZXMuaW8vc2VydmljZWZjY291bnQvc2Vyd
mljZS1hY2NvdW50LnVpZCI6IjY2YmRjNjM2LTJlYzMtNDJhZC00OGE1LWFhYzFlZjZlOTZlNSIsInN1YiI6InN5c3RlbTpzZXJ2aWNlYWNjb3
VudDpuZXB0dW5lOm5lcHRlbnUtc2EtZjIifQ.VYgboM4CTd0pdCJ78wjUwmtalh-2vsKjANyPsh-6guEwOtWEq5Fbw5ZHPtvAdrLlPzpOHEbA
e4eUM95BRGWbYIdwjuN95J0D4RNFkVUt48twoakRV7h-aPuwsQXhHZZzy4yimFHG9Ufmsk5Yr4RVcG6n137y-FH08K8zZjIPAsKDqNBQtXg-
lZvwVMi6viIhrrzAQs0MBOV820JYgy2o-WQVc0UUanCf94Y3gT0YTiQVvczYMs6nz9ut-
XgwitrBY6Tj9BgPprA9k_j5qEx_LUUZUpPAiEN7OzdkJsI8ctth101ypI1AeFr43t6ALyrQoBM39abDfq3FksR-oc_WMw
```

Question 6 | ReadinessProbe

Solve this question on instance: `ssh ckad5601`

Create a single *Pod* named `pod6` in *Namespace* `default` of image `busybox:1.31.0`. The *Pod* should have a readiness-probe executing `cat /tmp/ready`. It should initially wait 5 and periodically wait 10 seconds. This will set the container ready only if the file `/tmp/ready` exists.

The *Pod* should run the command `touch /tmp/ready && sleep 1d`, which will create the necessary file to be ready and then idles. Create the *Pod* and confirm it starts.

Answer:

```
k run pod6 --image=busybox:1.31.0 --dry-run=client -oyaml --command -- sh -c "touch /tmp/ready && sleep 1d" > 6.yaml

vim 6.yaml
```

Search for a readiness-probe example on <https://kubernetes.io/docs>, then copy and alter the relevant section for the task:

```
# 6.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: pod6
  name: pod6
spec:
  containers:
  - command:
    - sh
    - -c
    - touch /tmp/ready && sleep 1d
    image: busybox:1.31.0
    name: pod6
    resources: {}
    readinessProbe:
      exec:
        command:
        - sh
        - -c
        - cat /tmp/ready
      initialDelaySeconds: 5
      periodSeconds: 10
    dnsPolicy: ClusterFirst
    restartPolicy: Always
status: {}
```

Then:

```
k -f 6.yaml create
```

Running `k get pod6` we should see the job being created and completed:

```
→ k get pod pod6
```

NAME	READY	STATUS	RESTARTS	AGE
pod6	0/1	ContainerCreating	0	2s

```
→ k get pod pod6
```

NAME	READY	STATUS	RESTARTS	AGE
pod6	0/1	Running	0	7s

```
→ k get pod pod6
```

NAME	READY	STATUS	RESTARTS	AGE
pod6	1/1	Running	0	15s

We see that the *Pod* is finally ready.

Question 7 | Pods, Namespaces

Solve this question on instance: `ssh ckad7326`

The board of Team Neptune decided to take over control of one e-commerce webserver from Team Saturn. The administrator who once setup this webserver is not part of the organisation any longer. All information you could get was that the e-commerce system is called `my-happy-shop`.

Search for the correct *Pod* in *Namespace* `saturn` and move it to *Namespace* `neptune`. It doesn't matter if you shut it down and spin it up again, it probably hasn't any customers anyways.

Answer:

Let's see all those *Pods*:

```
→ k -n saturn get pod
```

NAME	READY	STATUS	RESTARTS	AGE
webserver-sat-001	1/1	Running	0	111m
webserver-sat-002	1/1	Running	0	111m
webserver-sat-003	1/1	Running	0	111m
webserver-sat-004	1/1	Running	0	111m
webserver-sat-005	1/1	Running	0	111m
webserver-sat-006	1/1	Running	0	111m

The *Pod* names don't reveal any information. We assume the *Pod* we are searching has a *label* or *annotation* with the name `my-happy-shop`, so we search for it:

```
k -n saturn describe pod # describe all pods, then manually look for it
```

```
# or do some filtering like this
```

```
k -n saturn get pod -o yaml | grep my-happy-shop -A10
```

We see the webserver we're looking for is `webserver-sat-003`

```
k -n saturn get pod webserver-sat-003 -o yaml > 7_webserver-sat-003.yaml # export
vim 7_webserver-sat-003.yaml
```

Change the *Namespace* to `neptune`, also remove the `status:` section, the token `volume`, the token `volumeMount` and the `nodeName`, else the new *Pod* won't start. The final file could look as clean like this:

```
# 7_webserver-sat-003.yaml
apiVersion: v1
kind: Pod
metadata:
  annotations:
    description: this is the server for the E-Commerce System my-happy-shop
  labels:
    id: webserver-sat-003
    name: webserver-sat-003
    namespace: neptune # new namespace here
spec:
  containers:
    - image: nginx:1.16.1-alpine
      imagePullPolicy: IfNotPresent
      name: webserver-sat
      restartPolicy: Always
```

Then we execute:

```
k -n neptune create -f 7_webserver-sat-003.yaml
```

```
→ k -n neptune get pod | grep webserver
webserver-sat-003          1/1      Running          0          22s
```

It seems the server is running in *Namespace* `neptune`, so we can do:

```
k -n saturn delete pod webserver-sat-003 --force --grace-period=0
```

Let's confirm only one is running:

```
→ k get pod -A | grep webserver-sat-003
neptune          webserver-sat-003          1/1      Running          0          6s
```

This should list only one pod called `webserver-sat-003` in *Namespace* `neptune`, status running.

Question 8 | Deployment, Rollouts

Solve this question on instance: `ssh ckad7326`

There is an existing *Deployment* named `api-new-c32` in *Namespace* `neptune`. A developer did make an update to the *Deployment* but the updated version never came online. Check the *Deployment* history and find a revision that works, then rollback to it. Could you tell Team Neptune what the error was so it doesn't happen again?

Answer:

```
k -n neptune get deploy # overview
k -n neptune rollout -h
k -n neptune rollout history -h
```

```
→ k -n neptune rollout history deploy api-new-c32
deployment.extensions/api-new-c32
REVISION  CHANGE-CAUSE
1          <none>
2          kubectl edit deployment api-new-c32 --namespace=neptune
3          kubectl edit deployment api-new-c32 --namespace=neptune
4          kubectl edit deployment api-new-c32 --namespace=neptune
5          kubectl edit deployment api-new-c32 --namespace=neptune
```

We see 5 revisions, let's check *Pod* and *Deployment* status:

```
→ k -n neptune get deploy,pod | grep api-new-c32
deployment.extensions/api-new-c32      3/3      1          3          141m

pod/api-new-c32-65d998785d-jtmqq      1/1      Running    0          141m
pod/api-new-c32-686d6f6b65-mj2fp      1/1      Running    0          141m
pod/api-new-c32-6dd45bdb68-2p462      1/1      Running    0          141m
pod/api-new-c32-7d64747c87-zh648      0/1      ImagePullBackOff 0          141m
```

Let's check the pod for errors:

```
→ k -n neptune describe pod api-new-c32-7d64747c87-zh648 | grep -i error
...   Error: ImagePullBackOff
```

```
→ k -n neptune describe pod api-new-c32-7d64747c87-zh648 | grep -i image
Image:          nginx:1.16.3
Image ID:
Reason:         ImagePullBackOff
Warning Failed 4m28s (x616 over 144m) kubelet, gke-s3ef67020-28c5-45f7--default-pool-248abd4f-s010
Error: ImagePullBackOff
```

Someone seems to have added a new image with a spelling mistake in the name `nginx:1.16.3`, that's the reason we can tell Team Neptune!

Now let's revert to the previous version:

```
k -n neptune rollout undo deploy api-new-c32
```

Does this one work?

```
→ k -n neptune get deploy api-new-c32
NAME          READY    UP-TO-DATE    AVAILABLE    AGE
api-new-c32   3/3      3             3            146m
```

Yes! All up-to-date and available.

Also a fast way to get an overview of the *ReplicaSets* of a *Deployment* and their images could be done with:

```
k -n neptune get rs -o wide | grep api-new-c32
```

Question 9 | Pod -> Deployment

Solve this question on instance: `ssh ckad9043`

In *Namespace* `pluto` there is single *Pod* named `holy-api`. It has been working okay for a while now but Team Pluto needs it to be more reliable.

Convert the *Pod* into a *Deployment* named `holy-api` with 3 replicas and delete the single *Pod* once done. The raw *Pod* template file is available at `/opt/course/9/holy-api-pod.yaml`.

In addition, the new *Deployment* should set `allowPrivilegeEscalation: false` and `privileged: false` for the security context on container level.

Please create the *Deployment* and save its yaml under `/opt/course/9/holy-api-deployment.yaml` on `ckad9043`.

Answer

There are multiple ways to do this, one is to copy an *Deployment* example from <https://kubernetes.io/docs> and then merge it with the existing *Pod* yaml. That's what we will do now:

```
cp /opt/course/9/holy-api-pod.yaml /opt/course/9/holy-api-deployment.yaml # make a copy!

vim /opt/course/9/holy-api-deployment.yaml
```

Now copy/use a *Deployment* example yaml and put the *Pod*'s **metadata:** and **spec:** into the *Deployment*'s **template:** section:

```
# /opt/course/9/holy-api-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: holy-api          # name stays the same
  namespace: pluto        # important
spec:
  replicas: 3              # 3 replicas
  selector:
    matchLabels:
      id: holy-api         # set the correct selector
  template:
    # => from here down it's the same as the pods metadata: and spec: sections
    metadata:
      labels:
        id: holy-api
        name: holy-api
    spec:
      containers:
        - env:
            - name: CACHE_KEY_1
              value: b&MTCi0=[T66RXm!jO@
            - name: CACHE_KEY_2
              value: PCAILGej5Ld@Q%{Q1=#
            - name: CACHE_KEY_3
              value: 2qz-]2OJlWDSTn_;RFQ
          image: nginx:1.17.3-alpine
          name: holy-api-container
```

```

securityContext:
  allowPrivilegeEscalation: false # add
  privileged: false               # add
volumeMounts:
- mountPath: /cache1
  name: cache-volume1
- mountPath: /cache2
  name: cache-volume2
- mountPath: /cache3
  name: cache-volume3
volumes:
- emptyDir: {}
  name: cache-volume1
- emptyDir: {}
  name: cache-volume2
- emptyDir: {}
  name: cache-volume3

```

To indent multiple lines using `vim` you should set the shiftwidth using `:set shiftwidth=2`. Then mark multiple lines using `Shift v` and the up/down keys.

To then indent the marked lines press `>` or `<` and to repeat the action press `.`

Next create the new *Deployment*:

```
k -f /opt/course/9/holy-api-deployment.yaml create
```

and confirm it's running:

```

→ k -n pluto get pod | grep holy
NAME                                READY   STATUS    RESTARTS   AGE
holy-api                            1/1     Running   0           19m
holy-api-5dbfdb4569-8qr5x           1/1     Running   0           30s
holy-api-5dbfdb4569-b5clh           1/1     Running   0           30s
holy-api-5dbfdb4569-rj2gz           1/1     Running   0           30s

```

Finally delete the single *Pod*:

```
k -n pluto delete pod holy-api --force --grace-period=0
```

```

→ k -n pluto get pod,deployment | grep holy
pod/holy-api-5dbfdb4569-8qr5x      1/1     Running   0           2m4s
pod/holy-api-5dbfdb4569-b5clh      1/1     Running   0           2m4s
pod/holy-api-5dbfdb4569-rj2gz      1/1     Running   0           2m4s

deployment.extensions/holy-api      3/3      3           3           2m4s

```

Question 10 | Service, Logs

Solve this question on instance: `ssh ckad9043`

Team Pluto needs a new cluster internal *Service*. Create a ClusterIP *Service* named `project-plt-6cc-svc` in *Namespace* `pluto`.

This *Service* should expose a single *Pod* named `project-plt-6cc-api` of image `nginx:1.17.3-alpine`, create that *Pod* as well. The *Pod* should be identified by label `project: plt-6cc-api`. The *Service* should use tcp port redirection of `3333:80`.

Finally use for example `curl` from a temporary `nginx:alpine` *Pod* to get the response from the *Service*. Write the response into `/opt/course/10/service_test.html` on `ckad9043`. Also check if the logs of *Pod* `project-plt-6cc-api` show the request and write those into `/opt/course/10/service_test.log` on `ckad9043`.

Answer

```
k -n pluto run project-plt-6cc-api --image=nginx:1.17.3-alpine --labels project=plt-6cc-api
```

This will create the requested *Pod*. In yaml it would look like this:

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    project: plt-6cc-api
  name: project-plt-6cc-api
spec:
  containers:
  - image: nginx:1.17.3-alpine
    name: project-plt-6cc-api
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

Next we create the service:

```
k -n pluto expose pod -h # help

k -n pluto expose pod project-plt-6cc-api --name project-plt-6cc-svc --port 3333 --target-port 80
```

Expose will create a yaml where everything is already set for our case and no need to change anything:

```
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    project: plt-6cc-api
  name: project-plt-6cc-svc    # good
  namespace: pluto           # great
spec:
  ports:
  - port: 3333                # awesome
    protocol: TCP
    targetPort: 80            # nice
  selector:
    project: plt-6cc-api      # beautiful
status:
  loadBalancer: {}
```

We could also use `create service` but then we would need to change the yaml afterwards:


```
k -n pluto create service -h # help
k -n pluto create service clusterip -h #help
k -n pluto create service clusterip project-plt-6cc-svc --tcp 3333:80 --dry-run=client -oyaml
# now we would need to set the correct selector labels
```

Check the *Service* is running:

```
→ k -n pluto get pod,svc | grep 6cc
```

pod/project-plt-6cc-api	1/1	Running	0	9m42s
-------------------------	-----	---------	---	-------

service/project-plt-6cc-svc	ClusterIP	10.31.241.234	<none>	3333/TCP	2m24s
-----------------------------	-----------	---------------	--------	----------	-------

Does the *Service* has one *Endpoint*?

```
→ k -n pluto describe svc project-plt-6cc-svc
```

Name: project-plt-6cc-svc
Namespace: pluto
Labels: project=plt-6cc-api
Annotations: <none>
Selector: project=plt-6cc-api
Type: ClusterIP
IP: 10.3.244.240
Port: <unset> 3333/TCP
TargetPort: 80/TCP
Endpoints: 10.28.2.32:80
Session Affinity: None
Events: <none>

Or even shorter:

```
→ k -n pluto get ep
```

NAME	ENDPOINTS	AGE
project-plt-6cc-svc	10.28.2.32:80	84m

Yes, endpoint there! Finally we check the connection using a temporary *Pod*:

```
→ k run tmp --restart=Never --rm --image=nginx:alpine -i -- curl http://project-plt-6cc-svc.pluto:3333
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current	
			Dload	Upload	Total	Spent	Left	Speed
100	612	100	612	0	0	32210	0	--:--:-- --:--:-- --:--:-- 32210

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
...
```

Great! Notice that we use the Kubernetes *Namespace* dns resolving (`project-plt-6cc-svc.pluto`) here. We could only use the

Service name if we would also spin up the temporary Pod in *Namespace* `pluto` .

And now really finally copy or pipe the html content into `/opt/course/10/service_test.html` .

```
# /opt/course/10/service_test.html
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
    body {
        width: 35em;
        margin: 0 auto;
        font-family: Tahoma, Verdana, Arial, sans-serif;
    }
...

```

Also the requested logs:

```
k -n pluto logs project-plt-6cc-api > /opt/course/10/service_test.log

```


```
# /opt/course/10/service_test.log
10.44.0.0 - - [22/Jan/2021:23:19:55 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.69.1" "-"

```

Question 11 | Working with Containers

Solve this question on instance: `ssh ckad9043`

There are files to build a container image located at `/opt/course/11/image` on `ckad9043` . The container will run a Golang application which outputs information to stdout. You're asked to perform the following tasks:

 Run all Docker and Podman commands as user root. Use `sudo docker` and `sudo podman` or become root with `sudo -i`

1. Change the Dockerfile: set ENV variable `SUN_CIPHER_ID` to hardcoded value `5b9c1065-e39d-4a43-a04a-e59bcea3e03f`
2. Build the image using `sudo docker`, tag it `registry.killer.sh:5000/sun-cipher:v1-docker` and push it to the registry
3. Build the image using `sudo podman`, tag it `registry.killer.sh:5000/sun-cipher:v1-podman` and push it to the registry
4. Run a container using `sudo podman`, which keeps running detached in the background, named `sun-cipher` using image `registry.killer.sh:5000/sun-cipher:v1-podman`
5. Write the logs your container `sun-cipher` produces into `/opt/course/11/logs` on `ckad9043`

Answer

Dockerfile: list of commands from which an *Image* can be build

Image: binary file which includes all data/requirements to be run as a *Container*

Container: running instance of an *Image*

Registry: place where we can push/pull *Images* to/from

Step 1

We should probably create a backup:

```
→ cp /opt/course/11/image/Dockerfile /opt/course/11/image/Dockerfile_bak
```

First we need to change the `/opt/course/11/logs/Dockerfile` to:

```
# build container stage 1
FROM docker.io/library/golang:1.15.15-alpine3.14
WORKDIR /src
COPY . .
RUN CGO_ENABLED=0 GOOS=linux go build -a -installsuffix cgo -o bin/app .

# app container stage 2
FROM docker.io/library/alpine:3.12.4
COPY --from=0 /src/bin/app app
# CHANGE NEXT LINE
ENV SUN_CIPHER_ID=5b9c1065-e39d-4a43-a04a-e59bcea3e03f
CMD ["/app"]
```

Step 2

Then we build the image using Docker:

```
→ cd /opt/course/11/image

→ sudo docker build -t registry.killer.sh:5000/sun-cipher:v1-docker .
...
Successfully built 409fde3c5bf9
Successfully tagged registry.killer.sh:5000/sun-cipher:v1-docker

→ sudo docker image ls
REPOSITORY                                TAG          IMAGE ID          CREATED          SIZE
registry.killer.sh:5000/sun-cipher        v1-docker    409fde3c5bf9      24 seconds ago  7.76MB
...

→ sudo docker push registry.killer.sh:5000/sun-cipher:v1-docker
The push refers to repository [registry.killer.sh:5000/sun-cipher]
c947fb5eba52: Pushed
33e8713114f8: Pushed
latest: digest: sha256:d216b4136a5b232b738698e826e7d12fccba9921d163b63777be23572250f23d size: 739
```

There we go, built and pushed.

Step 3

Next we build the image using Podman. Here it's only required to create one tag. The usage of Podman is very similar (for most cases even identical) to Docker:

```

→ cd /opt/course/11/image

→ sudo podman build -t registry.killer.sh:5000/sun-cipher:v1-podman .
...
--> 38adc53bd92
Successfully tagged registry.killer.sh:5000/sun-cipher:v1-podman
38adc53bd92881d91981c4b537f4f1b64f8de1de1b32eacc8479883170cee537

→ sudo podman image ls
REPOSITORY                                TAG          IMAGE ID          CREATED          SIZE
registry.killer.sh:5000/sun-cipher        v1-podman    38adc53bd928      2 minutes ago   8.03 MB
...

→ sudo podman push registry.killer.sh:5000/sun-cipher:v1-podman
Getting image source signatures
Copying blob 4d0d60db9eb6 done
Copying blob 33e8713114f8 done
Copying config bfala225f8 done
Writing manifest to image destination
Storing signatures

```

Built and pushed using Podman.

Step 4

We'll create a container from the perviously created image, using Podman, which keeps running in the background:

```

→ sudo podman run -d --name sun-cipher registry.killer.sh:5000/sun-cipher:v1-podman
f8199cba792f9fd2d1bd4decc9b7a9c0acfb975d95eda35f5f583c9efbf95589

```

Step 5

Finally we need to collect some information into files:

```

→ sudo podman logs sun-cipher
2077/03/13 06:50:34 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 8081
2077/03/13 06:50:34 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 7887
2077/03/13 06:50:34 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 1847
2077/03/13 06:50:34 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 4059
2077/03/13 06:50:34 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 2081
2077/03/13 06:50:34 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 1318
2077/03/13 06:50:34 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 4425
2077/03/13 06:50:34 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 2540
2077/03/13 06:50:34 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 456
2077/03/13 06:50:34 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 3300
2077/03/13 06:50:34 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 694
2077/03/13 06:50:34 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 8511
2077/03/13 06:50:44 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 8162
2077/03/13 06:50:54 random number for 5b9c1065-e39d-4a43-a04a-e59bcea3e03f is 5089

→ sudo podman logs sun-cipher > /opt/course/11/logs

```

This is looking not too bad at all. Our container skills are back in town!

Question 12 | Storage, PV, PVC, Pod volume

Solve this question on instance: `ssh ckad5601`

Create a new *PersistentVolume* named `earth-project-earthflower-pv`. It should have a capacity of *2Gi*, accessMode *ReadWriteOnce*, hostPath `/Volumes/Data` and no storageClassName defined.

Next create a new *PersistentVolumeClaim* in *Namespace* `earth` named `earth-project-earthflower-pvc`. It should request *2Gi* storage, accessMode *ReadWriteOnce* and should not define a storageClassName. The *PVC* should bound to the *PV* correctly.

Finally create a new *Deployment* `project-earthflower` in *Namespace* `earth` which mounts that volume at `/tmp/project-data`. The *Pods* of that *Deployment* should be of image `httpd:2.4.41-alpine`.

Answer

```
vim 12_pv.yaml
```

Find an example from <https://kubernetes.io/docs> and alter it:

```
# 12_pv.yaml
kind: PersistentVolume
apiVersion: v1
metadata:
  name: earth-project-earthflower-pv
spec:
  capacity:
    storage: 2Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/Volumes/Data"
```

Then create it:

```
k -f 12_pv.yaml create
```

Next the *PersistentVolumeClaim*:

```
vim 12_pvc.yaml
```

Find an example from <https://kubernetes.io/docs> and alter it:

```
# 12_pvc.yaml
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: earth-project-earthflower-pvc
  namespace: earth
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 2Gi
```

Then create:

```
k -f 12_pvc.yaml create
```

And check that both have the status Bound:

```
→ k -n earth get pv,pvc
```

NAME	CAPACITY	ACCESS MODES	...	STATUS	CLAIM
persistentvolume/...earthflower-pv	2Gi	RWO	...	Bound	...er-pvc

NAME	STATUS	VOLUME	CAPACITY
persistentvolumeclaim/...earthflower-pvc	Bound	earth-project-earthflower-pv	2Gi

Next we create a *Deployment* and mount that volume:

```
k -n earth create deploy project-earthflower --image=httpd:2.4.41-alpine --dry-run=client -oyaml >
12_dep.yaml

vim 12_dep.yaml
```

Alter the yaml to mount the volume:

```
# 12_dep.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: project-earthflower
  name: project-earthflower
  namespace: earth
spec:
  replicas: 1
  selector:
    matchLabels:
      app: project-earthflower
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: project-earthflower
    spec:
      volumes:
        # add
      - name: data
        # add
        persistentVolumeClaim:
          # add
          claimName: earth-project-earthflower-pvc
        # add
      containers:
```

```
- image: httpd:2.4.41-alpine
  name: container
  volumeMounts:
    # add
  - name: data
    # add
    mountPath: /tmp/project-data
    # add
```

```
k -f 12_dep.yaml create
```

We can confirm it's mounting correctly:

```
→ k -n earth describe pod project-earthflower-d6887f7c5-pn5wv | grep -A2 Mounts:
Mounts:
  /tmp/project-data from data (rw) # there it is
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-n2sjj (ro)
```

Question 13 | Storage, StorageClass, PVC

Solve this question on instance: `ssh ckad9043`

Team Moonpie, which has the *Namespace* `moon`, needs more storage. Create a new *PersistentVolumeClaim* named `moon-pvc-126` in that namespace. This claim should use a new *StorageClass* `moon-retain` with the *provisioner* set to `moon-retainer` and the *reclaimPolicy* set to *Retain*. The claim should request storage of *3Gi*, an *accessMode* of *ReadWriteOnce* and should use the new *StorageClass*.

The provisioner `moon-retainer` will be created by another team, so it's expected that the *PVC* will not boot yet. Confirm this by writing the event message from the *PVC* into file `/opt/course/13/pvc-126-reason` on `ckad9043`.

Answer

```
vim 13_sc.yaml
```

Head to <https://kubernetes.io/docs>, search for "storageclass" and alter the example code to this:

```
# 13_sc.yaml
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: moon-retain
provisioner: moon-retainer
reclaimPolicy: Retain
```

```
k create -f 13_sc.yaml
```

Now the same for the *PersistentVolumeClaim*, head to the docs, copy an example and transform it into:

```
vim 13_pvc.yaml
```

```
# 13_pvc.yaml
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: moon-pvc-126          # name as requested
  namespace: moon            # important
spec:
  accessModes:
    - ReadWriteOnce          # RWO
  resources:
    requests:
      storage: 3Gi           # size
  storageClassName: moon-retain # uses our new storage class
```

```
k -f 13_pvc.yaml create
```

Next we check the status of the *PVC* :

```
→ k -n moon get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
moon-pvc-126	Pending				moon-retain	2m57s

```
→ k -n moon describe pvc moon-pvc-126
Name:          moon-pvc-126
...
Status:        Pending
...
Events:
  Type      Reason              Age             From              Message
  ----      -
Normal      ExternalProvisioning  4s (x19 over 4m28s)  persistentvolume-controller  Waiting for a volume to be created either by the external provisioner 'moon-retainer' or manually by the system administrator. If volume creation is delayed, please verify that the provisioner is running and correctly registered.
```

This confirms that the *PVC* waits for the provisioner `moon-retainer` to be created. Finally we copy or write the event message into the requested location:

```
# /opt/course/13/pvc-126-reason
Waiting for a volume to be created either by the external provisioner 'moon-retainer' or manually by the system administrator. If volume creation is delayed, please verify that the provisioner is running and correctly registered.
```

Question 14 | Secret, Secret-Volume, Secret-Env

Solve this question on instance: `ssh ckad9043`

You need to make changes on an existing *Pod* in *Namespace* `moon` called `secret-handler`. Create a new *Secret* `secret1` which contains `user=test` and `pass=pwd`. The *Secret*'s content should be available in *Pod* `secret-handler` as environment variables `SECRET1_USER` and `SECRET1_PASS`. The yaml for *Pod* `secret-handler` is available at `/opt/course/14/secret-handler.yaml`.

There is existing yaml for another *Secret* at `/opt/course/14/secret2.yaml`, create this *Secret* and mount it inside the same *Pod* at

/tmp/secret2. Your changes should be saved under /opt/course/14/secret-handler-new.yaml on ckad9043. Both *Secrets* should only be available in *Namespace* moon.

Answer

```
k -n moon get pod # show pods
k -n moon create secret -h # help
k -n moon create secret generic -h # help
k -n moon create secret generic secret1 --from-literal user=test --from-literal pass=pwd
```

The last command would generate this yaml:

```
apiVersion: v1
data:
  pass: cHdk
  user: dGVzdA==
kind: Secret
metadata:
  creationTimestamp: null
  name: secret1
  namespace: moon
```

Next we create the second *Secret* from the given location, making sure it'll be created in *Namespace* moon:

```
k -n moon -f /opt/course/14/secret2.yaml create
```

```
→ k -n moon get secret
```

NAME	TYPE	DATA	AGE
default-token-rvzcf	kubernetes.io/service-account-token	3	66m
secret1	Opaque	2	4m3s
secret2	Opaque	1	8s

We will now edit the *Pod* yaml:

```
cp /opt/course/14/secret-handler.yaml /opt/course/14/secret-handler-new.yaml
vim /opt/course/14/secret-handler-new.yaml
```

Add the following to the yaml:

```
# /opt/course/14/secret-handler-new.yaml
apiVersion: v1
kind: Pod
metadata:
  labels:
    id: secret-handler
    uuid: 1428721e-8d1c-4c09-b5d6-afd79200c56a
    red_ident: 9cf7a7c0-fdb2-4c35-9c13-c2a0bb52b4a9
    type: automatic
  name: secret-handler
  namespace: moon
spec:
  volumes:
  - name: cache-volume1
    emptyDir: {}
  - name: cache-volume2
    emptyDir: {}
  - name: cache-volume3
    emptyDir: {}
```

```

- name: secret2-volume                # add
  secret:                             # add
    secretName: secret2              # add
containers:
- name: secret-handler
  image: bash:5.0.11
  args: ['bash', '-c', 'sleep 2d']
  volumeMounts:
  - mountPath: /cache1
    name: cache-volume1
  - mountPath: /cache2
    name: cache-volume2
  - mountPath: /cache3
    name: cache-volume3
  - name: secret2-volume              # add
    mountPath: /tmp/secret2          # add
  env:
  - name: SECRET_KEY_1
    value: ">8$kH#kj..i8}HImQd{"
  - name: SECRET_KEY_2
    value: "IO=a4L/XkRdvN8jM=Y+"
  - name: SECRET_KEY_3
    value: "-7PA0_Z]>{pwa43r)___"
  - name: SECRET1_USER                # add
    valueFrom:                       # add
      secretKeyRef:                 # add
        name: secret1              # add
        key: user                   # add
  - name: SECRET1_PASS                # add
    valueFrom:                       # add
      secretKeyRef:                 # add
        name: secret1              # add
        key: pass                   # add

```

There is also the possibility to import all keys from a *Secret* as env variables at once, though the env variable names will then be the same as in the *Secret*, which doesn't work for the requirements here:

```

containers:
- name: secret-handler
...
  envFrom:
  - secretRef:                      # also works for configMapRef
    name: secret1

```

Then we apply the changes:

```

k -f /opt/course/14/secret-handler.yaml delete --force --grace-period=0
k -f /opt/course/14/secret-handler-new.yaml create

```

Instead of running delete and create we can also use recreate:

```

k -f /opt/course/14/secret-handler-new.yaml replace --force --grace-period=0

```

It was not requested directly, but you should always confirm it's working:

```
→ k -n moon exec secret-handler -- env | grep SECRET1
SECRET1_USER=test
SECRET1_PASS=pwd

→ k -n moon exec secret-handler -- find /tmp/secret2
/tmp/secret2
/tmp/secret2/..data
/tmp/secret2/key
/tmp/secret2/..2019_09_11_09_03_08.147048594
/tmp/secret2/..2019_09_11_09_03_08.147048594/key

→ k -n moon exec secret-handler -- cat /tmp/secret2/key
12345678
```

Question 15 | ConfigMap, Configmap-Volume

Solve this question on instance: `ssh ckad9043`

Team Moonpie has a nginx server *Deployment* called `web-moon` in *Namespace* `moon`. Someone started configuring it but it was never completed. To complete please create a *ConfigMap* called `configmap-web-moon-html` containing the content of file `/opt/course/15/web-moon.html` under the data key-name `index.html`.

The *Deployment* `web-moon` is already configured to work with this *ConfigMap* and serve its content. Test the nginx configuration for example using `curl` from a temporary `nginx:alpine` *Pod*.

Answer

Let's check the existing *Pods*:

```
→ k -n moon get pod
```

NAME	READY	STATUS	RESTARTS	AGE
secret-handler	1/1	Running	0	55m
web-moon-847496c686-2rzj4	0/1	ContainerCreating	0	33s
web-moon-847496c686-9nwwj	0/1	ContainerCreating	0	33s
web-moon-847496c686-cxdbx	0/1	ContainerCreating	0	33s
web-moon-847496c686-hvqlw	0/1	ContainerCreating	0	33s
web-moon-847496c686-tj7ct	0/1	ContainerCreating	0	33s

```
→ k -n moon describe pod web-moon-847496c686-2rzj4
...
Warning  FailedMount  31s (x7 over 63s)  kubelet, gke-test-default-pool-ce83a51a-p6s4  MountVolume.SetUp
failed for volume "html-volume" : configmaps "configmap-web-moon-html" not found
```

Good so far, now let's create the missing *ConfigMap*:

```
k -n moon create configmap -h # help

k -n moon create configmap configmap-web-moon-html --from-file=index.html=/opt/course/15/web-moon.html #
important to set the index.html key
```

This should create a *ConfigMap* with yaml like:

```
apiVersion: v1
data:
  index.html: |      # notice the key index.html, this will be the filename when mounted
    <!DOCTYPE html>
    <html lang="en">
    <head>
      <meta charset="UTF-8">
      <title>Web Moon Webpage</title>
    </head>
    <body>
      This is some great content.
    </body>
    </html>
kind: ConfigMap
metadata:
  creationTimestamp: null
  name: configmap-web-moon-html
  namespace: moon
```

After waiting a bit or deleting/recreating (`k -n moon rollout restart deploy web-moon`) the *Pods* we should see:

```
→ k -n moon get pod
```

NAME	READY	STATUS	RESTARTS	AGE
secret-handler	1/1	Running	0	59m
web-moon-847496c686-2rzj4	1/1	Running	0	4m28s
web-moon-847496c686-9nwwj	1/1	Running	0	4m28s
web-moon-847496c686-cxdbx	1/1	Running	0	4m28s
web-moon-847496c686-hvqlw	1/1	Running	0	4m28s
web-moon-847496c686-tj7ct	1/1	Running	0	4m28s

Looking much better. Finally we check if the nginx returns the correct content:

```
k -n moon get pod -o wide # get pod cluster IPs
```

Then use one IP to test the configuration:

```
→ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl 10.44.0.78
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload	Upload	Total	Spent	Left
100	161	100	161	0	0	80500	0

```
--:--:-- --:--:-- --:--:-- 157k
```

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>Web Moon Webpage</title>
</head>
<body>
  This is some great content.
</body>
```

For debugging or further checks we could find out more about the *Pods* volume mounts:

```
→ k -n moon describe pod web-moon-c77655cc-dc8v4 | grep -A2 Mounts:
Mounts:
  /usr/share/nginx/html from html-volume (rw)
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-rvzcf (ro)
```

And check the mounted folder content:

```
→ k -n moon exec web-moon-c77655cc-dc8v4 find /usr/share/nginx/html
/usr/share/nginx/html
/usr/share/nginx/html/..2019_09_11_10_05_56.336284411
/usr/share/nginx/html/..2019_09_11_10_05_56.336284411/index.html
/usr/share/nginx/html/..data
/usr/share/nginx/html/index.html
```

Here it was important that the file will have the name `index.html` and not the original one `web-moon.html` which is controlled through the *ConfigMap* data key.

Question 16 | Logging sidecar

Solve this question on instance: `ssh ckad7326`

The Tech Lead of Mercury2D decided it's time for more logging, to finally fight all these missing data incidents. There is an existing container named `cleaner-con` in *Deployment* `cleaner` in *Namespace* `mercury`. This container mounts a volume and writes logs into a file called `cleaner.log`.

The yaml for the existing *Deployment* is available at `/opt/course/16/cleaner.yaml`. Persist your changes at `/opt/course/16/cleaner-new.yaml` on `ckad7326` but also make sure the *Deployment* is running.

Create a sidecar container named `logger-con`, image `busybox:1.31.0`, which mounts the same volume and writes the content of `cleaner.log` to stdout, you can use the `tail -f` command for this. This way it can be picked up by `kubectl logs`.

Check if the logs of the new container reveal something about the missing data incidents.

Answer

Sidecar containers in K8s are `initContainers` with `restartPolicy: Always`. Search for "Sidecar Containers" in the K8s Docs to familiarise yourself if necessary.

```
cp /opt/course/16/cleaner.yaml /opt/course/16/cleaner-new.yaml
vim /opt/course/16/cleaner-new.yaml
```

Add a sidecar container which outputs the log file to stdout:

```
# /opt/course/16/cleaner-new.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  name: cleaner
  namespace: mercury
spec:
  replicas: 2
  selector:
    matchLabels:
      id: cleaner
  template:
    metadata:
```

```

labels:
  id: cleaner
spec:
  volumes:
  - name: logs
    emptyDir: {}
  initContainers:
  - name: init
    image: bash:5.0.11
    command: ['bash', '-c', 'echo init > /var/log/cleaner/cleaner.log']
    volumeMounts:
    - name: logs
      mountPath: /var/log/cleaner
  - name: logger-con                                     # add
    image: busybox:1.31.0                                # add
    restartPolicy: Always                                # add
    command: ["sh", "-c", "tail -f /var/log/cleaner/cleaner.log"] # add
    volumeMounts:                                         # add
    - name: logs                                           # add
      mountPath: /var/log/cleaner                         # add
  containers:
  - name: cleaner-con
    image: bash:5.0.11
    args: ['bash', '-c', 'while true; do echo `date`: "remove random file" >> /var/log/cleaner/
cleaner.log; sleep 1; done']
    volumeMounts:
    - name: logs
      mountPath: /var/log/cleaner

```

In earlier K8s versions it was necessary to define sidecar containers as additional application containers under `containers:` like this:

```

# LEGACY example of defining sidecar containers in earlier K8s versions
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  name: cleaner
  namespace: mercury
spec:
...
  template:
...
    spec:
...
      initContainers:
      - name: init
        image: bash:5.0.11
...
      containers:
      - name: cleaner-con
        image: bash:5.0.11
...
      - name: logger-con                                     # LEGACY example
        image: busybox:1.31.0                                # LEGACY example
        command: ["sh", "-c", "tail -f /var/log/cleaner/cleaner.log"] # LEGACY example
        volumeMounts:                                         # LEGACY example
        - name: logs                                           # LEGACY example
          mountPath: /var/log/cleaner                         # LEGACY example

```

Then apply the changes and check the logs of the sidecar:

```
k -f /opt/course/16/cleaner-new.yaml apply
```

This will cause a deployment rollout of which we can get more details:

```
k -n mercury rollout history deploy cleaner
k -n mercury rollout history deploy cleaner --revision 1
k -n mercury rollout history deploy cleaner --revision 2
```

Check *Pod* statuses:

```
→ k -n mercury get pod
```

NAME	READY	STATUS	RESTARTS	AGE
cleaner-86b7758668-9pw6t	2/2	Running	0	6s
cleaner-86b7758668-qgh4v	0/2	Init:0/1	0	1s


```
→ k -n mercury get pod
```

NAME	READY	STATUS	RESTARTS	AGE
cleaner-86b7758668-9pw6t	2/2	Running	0	14s
cleaner-86b7758668-qgh4v	2/2	Running	0	9s

Finally check the logs of the logging sidecar container:

```
→ k -n mercury logs cleaner-576967576c-cqtgx -c logger-con
init
Wed Sep 11 10:45:44 UTC 2099: remove random file
Wed Sep 11 10:45:45 UTC 2099: remove random file
...
```

Mystery solved, something is removing files at random ;) It's important to understand how containers can communicate with each other using volumes.

Question 17 | InitContainer

Solve this question on instance: `ssh ckad5601`

Last lunch you told your coworker from department Mars Inc how amazing *InitContainers* are. Now he would like to see one in action. There is a *Deployment* yaml at `/opt/course/17/test-init-container.yaml`. This *Deployment* spins up a single *Pod* of image `nginx:1.17.3-alpine` and serves files from a mounted volume, which is empty right now.

Create an *InitContainer* named `init-con` which also mounts that volume and creates a file `index.html` with content `check this out!` in the root of the mounted volume. For this test we ignore that it doesn't contain valid html.

The *InitContainer* should be using image `busybox:1.31.0`. Test your implementation for example using `curl` from a temporary `nginx:alpine` *Pod*.

Answer

```
cp /opt/course/17/test-init-container.yaml ~/17_test-init-container.yaml

vim 17_test-init-container.yaml
```

Add the *InitContainer*:

```
# 17_test-init-container.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: test-init-container
  namespace: mars
spec:
  replicas: 1
  selector:
    matchLabels:
      id: test-init-container
  template:
    metadata:
      labels:
        id: test-init-container
    spec:
      volumes:
      - name: web-content
        emptyDir: {}
      initContainers:
        - name: init-container # initContainer start
          image: busybox:1.31.0
          command: ['sh', '-c', 'echo "check this out!" > /tmp/web-content/index.html']
          volumeMounts:
            - name: web-content
              mountPath: /tmp/web-content # initContainer end
      containers:
      - image: nginx:1.17.3-alpine
        name: nginx
        volumeMounts:
          - name: web-content
            mountPath: /usr/share/nginx/html
        ports:
          - containerPort: 80
```

Then we create the *Deployment*:

```
k -f 17_test-init-container.yaml create
```

Finally we test the configuration:

```
k -n mars get pod -o wide # to get the cluster IP
```

```
→ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl 10.0.0.67
% Total      % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left  Speed

check this out!
```

Beautiful.

Question 18 | Service misconfiguration

Solve this question on instance: `ssh ckad5601`

There seems to be an issue in *Namespace* `mars` where the ClusterIP service `manager-api-svc` should make the *Pods* of *Deployment* `manager-api-deployment` available inside the cluster.

You can test this with `curl manager-api-svc.mars:4444` from a temporary `nginx:alpine` *Pod*. Check for the misconfiguration and apply a fix.

Answer

First let's get an overview:

```
→ k -n mars get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/manager-api-deployment-dbcc6657d-bg2hh	1/1	Running	0	98m
pod/manager-api-deployment-dbcc6657d-f5fv4	1/1	Running	0	98m
pod/manager-api-deployment-dbcc6657d-httjv	1/1	Running	0	98m
pod/manager-api-deployment-dbcc6657d-k98xn	1/1	Running	0	98m
pod/test-init-container-5db7c99857-htx6b	1/1	Running	0	2m19s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/manager-api-svc	ClusterIP	10.15.241.159	<none>	4444/TCP	99m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/manager-api-deployment	4/4	4	4	98m
deployment.apps/test-init-container	1/1	1	1	2m19s
...				

Everything seems to be running, but we can't seem to get a connection:

```
→ k -n mars run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 manager-api-svc:4444
If you don't see a command prompt, try pressing enter.
 0      0      0      0      0      0      0      0 --:--:--  0:00:01 --:--:--      0
curl: (28) Connection timed out after 1000 milliseconds
pod "tmp" deleted
pod mars/tmp terminated (Error)
```

Ok, let's try to connect to one pod directly:

```
k -n mars get pod -o wide # get cluster IP
```



```
→ k -n mars run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 10.0.1.14
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
<!DOCTYPE html>							
<html>							
<head>							
<title>Welcome to nginx!</title>							
...							

The *Pods* itself seem to work. Let's investigate the *Service* a bit:

```
→ k -n mars describe service manager-api-svc
Name:                manager-api-svc
Namespace:           mars
Labels:              app=manager-api-svc
...
Endpoints:           <none>
...
```

Endpoint inspection is also possible using:

```
k -n mars get ep
```

No endpoints - No good. We check the *Service* yaml:

```
k -n mars edit service manager-api-svc
```

```
# k -n mars edit service manager-api-svc
apiVersion: v1
kind: Service
metadata:
...
  labels:
    app: manager-api-svc
    name: manager-api-svc
    namespace: mars
...
spec:
  clusterIP: 10.3.244.121
  ports:
  - name: 4444-80
    port: 4444
    protocol: TCP
    targetPort: 80
  selector:
    #id: manager-api-deployment # wrong selector, needs to point to pod!
    id: manager-api-pod
  sessionAffinity: None
  type: ClusterIP
```

Though *Pods* are usually never created without a *Deployment* or *ReplicaSet*, *Services* always select for *Pods* directly. This gives great flexibility because *Pods* could be created through various customized ways. After saving the new selector we check the *Service* again for endpoints:

```
→ k -n mars get ep
NAME                ENDPOINTS
manager-api-svc     10.0.0.30:80,10.0.1.30:80,10.0.1.31:80 + 1 more... 41m
```

Endpoints - Good! Now we try connecting again:

```
→ k -n mars run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 manager-api-svc:4444
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left  Speed
100    612    100    612     0     0    99k      0  --:--:-- --:--:-- --:--:--   99k
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
...
```

And we fixed it. Good to know is how to be able to use Kubernetes DNS resolution from a different *Namespace*. Not necessary, but we could spin up the temporary *Pod* in default *Namespace*:

```
→ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 manager-api-svc:4444
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left     Speed

  0      0    0     0    0     0     0      0  --:--:-- --:--:-- --:--:--    0curl: (6) Could not resolve
host: manager-api-svc
pod "tmp" deleted
pod default/tmp terminated (Error)

→ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 manager-api-svc.mars:4444
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left     Speed

100    612   100    612    0     0 68000      0  --:--:-- --:--:-- --:--:-- 68000
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
```

Short `manager-api-svc.mars` or long `manager-api-svc.mars.svc.cluster.local` work.

Question 19 | Service ClusterIP->NodePort

Solve this question on instance: `ssh ckad5601`

In *Namespace* `jupiter` you'll find an apache *Deployment* (with one replica) named `jupiter-crew-deploy` and a ClusterIP *Service* called `jupiter-crew-svc` which exposes it. Change this service to a NodePort one to make it available on all nodes on port 30100.

Test the NodePort *Service* using the internal IP of all available nodes and the port 30100 using `curl`, you can reach the internal node IPs directly from your main terminal. On which nodes is the *Service* reachable? On which node is the *Pod* running?

Answer

First we get an overview:

```
→ k -n jupiter get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/jupiter-crew-deploy-8cdf99bc9-klwqt	1/1	Running	0	34m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/jupiter-crew-svc	ClusterIP	10.100.254.66	<none>	8080/TCP	34m
...					

(Optional) Next we check if the ClusterIP *Service* actually works:

```
→ k -n jupiter run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 jupiter-crew-svc:8080
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
100	45	100	45	0	0	5000	0
				--:--:--	--:--:--	--:--:--	5000

```
<html><body><h1>It works!</h1></body></html>
```

The *Service* is working great. Next we change the *Service* type to NodePort and set the port:

```
k -n jupiter edit service jupiter-crew-svc
```

```
# k -n jupiter edit service jupiter-crew-svc
apiVersion: v1
kind: Service
metadata:
  name: jupiter-crew-svc
  namespace: jupiter
...
spec:
  clusterIP: 10.3.245.70
  ports:
  - name: 8080-80
    port: 8080
    protocol: TCP
    targetPort: 80
    nodePort: 30100 # add the nodePort
  selector:
    id: jupiter-crew
  sessionAffinity: None
  #type: ClusterIP
  type: NodePort # change type
status:
  loadBalancer: {}
```

We check if the *Service* type was updated:

```
→ k -n jupiter get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
jupiter-crew-svc	NodePort	10.3.245.70	<none>	8080:30100/TCP	3m52s

(Optional) And we confirm that the service is still reachable internally:

```
→ k -n jupiter run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 jupiter-crew-svc:8080
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed

```
<html><body><h1>It works!</h1></body></html>
```

Nice. A NodePort *Service* kind of lies on top of a ClusterIP one, making the ClusterIP *Service* reachable on the Node IPs (internal and external). Next we get the *internal* IPs of all nodes to check the connectivity:

```
→ k get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	...
ckad5601	Ready	control-plane	18h	v1.33.1	192.168.100.11	...

We can test the connection using the node IP:

```
→ curl 192.168.100.11:30100
```

```
<html><body><h1>It works!</h1></body></html>
```

Here we only have one node in the cluster, but the *Service* would be reachable on all of them. Even if the *Pod* is just running on one specific node, the *Service* makes it available through port 30100 on the internal and external IP addresses of all nodes. This is at least the common/default behaviour but can depend on cluster configuration.

Question 20 | NetworkPolicy

Solve this question on instance: `ssh ckad7326`

In *Namespace* `venus` you'll find two *Deployments* named `api` and `frontend`. Both *Deployments* are exposed inside the cluster using *Services*. Create a *NetworkPolicy* named `np1` which restricts outgoing tcp connections from *Deployment* `frontend` and only allows those going to *Deployment* `api`. Make sure the *NetworkPolicy* still allows outgoing traffic on UDP/TCP ports 53 for DNS resolution.

Test using: `wget www.google.com` and `wget api:2222` from a *Pod* of *Deployment* `frontend`.

Answer

INFO: For learning NetworkPolicies check out <https://editor.cilium.io>. But you're not allowed to use it during the exam.

First we get an overview:

```
→ k -n venus get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/api-5979b95578-gktxp	1/1	Running	0	57s
pod/api-5979b95578-lhc15	1/1	Running	0	57s
pod/frontend-789cbdc677-c9v8h	1/1	Running	0	57s
pod/frontend-789cbdc677-npk2m	1/1	Running	0	57s
pod/frontend-789cbdc677-pl67g	1/1	Running	0	57s
pod/frontend-789cbdc677-rjt5r	1/1	Running	0	57s
pod/frontend-789cbdc677-xgf5n	1/1	Running	0	57s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/api	ClusterIP	10.3.255.137	<none>	2222/TCP	37s
service/frontend	ClusterIP	10.3.255.135	<none>	80/TCP	57s
...					

(Optional) This is not necessary but we could check if the *Services* are working inside the cluster:

```
→ k -n venus run tmp --restart=Never --rm -i --image=busybox -i -- wget -O- frontend:80
Connecting to frontend:80 (10.3.245.9:80)
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
...

→ k -n venus run tmp --restart=Never --rm --image=busybox -i -- wget -O- api:2222
Connecting to api:2222 (10.3.250.233:2222)
<html><body><h1>It works!</h1></body></html>
```

Then we use any `frontend` *Pod* and check if it can reach external names and the `api` *Service*:

```
→ k -n venus exec frontend-789cbdc677-c9v8h -- wget -O- www.google.com
Connecting to www.google.com (216.58.205.227:80)
- 100% |*****| 12955 0:00:00 ETA
<!doctype html><html itemscope="" itemtype="http://schema.org/WebPage" lang="en"><head>
...

→ k -n venus exec frontend-789cbdc677-c9v8h -- wget -O- api:2222
<html><body><h1>It works!</h1></body></html>
Connecting to api:2222 (10.3.255.137:2222)
- 100% |*****| 45 0:00:00 ETA
...
```

We see *Pods* of `frontend` can reach the `api` and external names.

```
vim 20_np1.yaml
```

Now we head to <https://kubernetes.io/docs>, search for *NetworkPolicy*, copy the example code and adjust it to:

```
# 20_np1.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: np1
  namespace: venus
spec:
  podSelector:
    matchLabels:
      id: frontend          # label of the pods this policy should be applied on
  policyTypes:
  - Egress                 # we only want to control egress
  egress:
  - to:                    # 1st egress rule
    - podSelector:         # allow egress only to pods with api label
      matchLabels:
        id: api
  - ports:                 # 2nd egress rule
    - port: 53              # allow DNS UDP
      protocol: UDP
    - port: 53              # allow DNS TCP
      protocol: TCP
```

Notice that we specify two egress rules in the yaml above. If we specify multiple egress rules then these are connected using a logical OR. So in the example above we do:

```
allow outgoing traffic if
  (destination pod has label id:api) OR ((port is 53 UDP) OR (port is 53 TCP))
```

Let's have a look at example code which wouldn't work in our case:

```
# this example does not work in our case
...
egress:
- to:                                # 1st AND ONLY egress rule
  - podSelector:                     # allow egress only to pods with api label
    matchLabels:
      id: api
  ports:                             # STILL THE SAME RULE but just an additional selector
  - port: 53                         # allow DNS UDP
    protocol: UDP
  - port: 53                         # allow DNS TCP
    protocol: TCP
```

In the yaml above we only specify one egress rule with two selectors. It can be translated into:

```
allow outgoing traffic if
  (destination pod has label id:api) AND ((port is 53 UDP) OR (port is 53 TCP))
```

Apply the correct policy:

```
k -f 20_npl.yaml create
```

And try again, external is not working any longer:

```
→ k -n venus exec frontend-789cbdc677-c9v8h -- wget -O- www.google.de
Connecting to www.google.de:2222 (216.58.207.67:80)
^C

→ k -n venus exec frontend-789cbdc677-c9v8h -- wget -O- -T 5 www.google.de:80
Connecting to www.google.com (172.217.203.104:80)
wget: download timed out
command terminated with exit code 1
```

Internal connection to `api` work as before:

```
→ k -n venus exec frontend-789cbdc677-c9v8h -- wget -O- api:2222
<html><body><h1>It works!</h1></body></html>
Connecting to api:2222 (10.3.255.137:2222)
- 100% |*****| 45 0:00:00 ETA
```

Question 21 | Requests and Limits, ServiceAccount

Solve this question on instance: `ssh ckad7326`

Team Neptune needs 3 *Pods* of image `httpd:2.4-alpine`, create a *Deployment* named `neptune-10ab` for this. The containers

should be named `neptune-pod-10ab`. Each container should have a memory request of *20Mi* and a memory limit of *50Mi*.

Team Neptune has its own *ServiceAccount* `neptune-sa-v2` under which the *Pods* should run. The *Deployment* should be in *Namespace* `neptune`.

Answer:

```
k -n neptune create deployment -h # help
k -n neptune create deploy -h # deploy is short for deployment

k -n neptune create deploy neptune-10ab --replicas=3 --image=httpd:2.4-alpine --dry-run=client -oyaml >
21.yaml

vim 21.yaml
```

Now make the required changes using vim:

```
# 21.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: neptune-10ab
  name: neptune-10ab
  namespace: neptune
spec:
  replicas: 3
  selector:
    matchLabels:
      app: neptune-10ab
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: neptune-10ab
    spec:
      serviceAccountName: neptune-sa-v2 # add
      containers:
      - image: httpd:2.4-alpine
        name: neptune-pod-10ab # change
        resources: # add
          limits: # add
            memory: 50Mi # add
          requests: # add
            memory: 20Mi # add
      status: {}
```

Then create the yaml:

```
k create -f 21.yaml # namespace already set in yaml
```

To verify all *Pods* are running we do:

```
→ k -n neptune get pod | grep neptune-10ab
```

neptune-10ab-7d4b8d45b-4nzj5	1/1	Running	0	57s
neptune-10ab-7d4b8d45b-lzwrp	1/1	Running	0	17s
neptune-10ab-7d4b8d45b-z5hcc	1/1	Running	0	17s

Question 22 | Labels, Annotations

Solve this question on instance: `ssh ckad9043`

Team Sunny needs to identify some of their *Pods* in namespace `sun`. They ask you to add a new label `protected: true` to all *Pods* with an existing label `type: worker` or `type: runner`. Also add an annotation `protected: do not delete this pod` to all *Pods* having the new label `protected: true`.

Answer

```
→ k -n sun get pod --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
0509649a	1/1	Running	0	25s	type=runner,type_old=messenger
0509649b	1/1	Running	0	24s	type=worker
1428721e	1/1	Running	0	23s	type=worker
1428721f	1/1	Running	0	22s	type=worker
43b9a	1/1	Running	0	22s	type=test
4c09	1/1	Running	0	21s	type=worker
4c35	1/1	Running	0	20s	type=worker
4fe4	1/1	Running	0	19s	type=worker
5555a	1/1	Running	0	19s	type=messenger
86cda	1/1	Running	0	18s	type=runner
8dlc	1/1	Running	0	17s	type=messenger
a004a	1/1	Running	0	16s	type=runner
a94128196	1/1	Running	0	15s	type=runner,type_old=messenger
afd79200c56a	1/1	Running	0	15s	type=worker
b667	1/1	Running	0	14s	type=worker
fdb2	1/1	Running	0	13s	type=worker

If we would only like to get pods with certain labels we can run:

```
k -n sun get pod -l type=runner # only pods with label runner
```

We can use this label filtering also when using other commands, like setting new labels:

```
k label -h # help
k -n sun label pod -l type=runner protected=true # run for label runner
k -n sun label pod -l type=worker protected=true # run for label worker
```

Or we could run:

```
k -n sun label pod -l "type in (worker,runner)" protected=true
```

Let's check the result:

```
→ k -n sun get pod --show-labels
```

NAME	...	AGE	LABELS
0509649a	...	56s	protected=true,type=runner,type_old=messenger
0509649b	...	55s	protected=true,type=worker

1428721e	...	54s	protected=true,type=worker
1428721f	...	53s	protected=true,type=worker
43b9a	...	53s	type=test
4c09	...	52s	protected=true,type=worker
4c35	...	51s	protected=true,type=worker
4fe4	...	50s	protected=true,type=worker
5555a	...	50s	type=messenger
86cda	...	49s	protected=true,type=runner
8dlc	...	48s	type=messenger
a004a	...	47s	protected=true,type=runner
a94128196	...	46s	protected=true,type=runner,type_old=messenger
afd79200c56a	...	46s	protected=true,type=worker
b667	...	45s	protected=true,type=worker
fdb2	...	44s	protected=true,type=worker

Looking good. Finally we set the annotation using the newly assigned label `protected: true`:

```
k -n sun annotate pod -l protected=true protected="do not delete this pod"
```

Not requested in the task but for your own control you could run:

```
k -n sun get pod -l protected=true -o yaml | grep -A 8 metadata:
```

CKAD Simulator Preview Kubernetes 1.33

<https://killer.sh>

This is a preview of the CKAD Simulator content. The full CKAD Simulator contains 22 different questions. These preview questions are in addition to the provided ones and can also be solved in the interactive environment.

Preview Question 1

Solve this question on instance: `ssh ckad9043`

In *Namespace* `pluto` there is a *Deployment* named `project-23-api`. It has been working okay for a while but Team Pluto needs it to be more reliable. Implement a liveness-probe which checks the container to be reachable on port 80. Initially the probe should wait *10*, periodically *15* seconds.

The original *Deployment* yaml is available at `/opt/course/p1/project-23-api.yaml`. Save your changes at `/opt/course/p1/project-23-api-new.yaml` and apply the changes.

Answer

First we get an overview:

```
→ k -n pluto get all -o wide
```

NAME	READY	STATUS	... IP	...
pod/holy-api	1/1	Running	... 10.12.0.26	...
pod/project-23-api-784857f54c-dx6h6	1/1	Running	... 10.12.2.15	...
pod/project-23-api-784857f54c-sj8df	1/1	Running	... 10.12.1.18	...
pod/project-23-api-784857f54c-t4xmh	1/1	Running	... 10.12.0.23	...

NAME	READY	UP-TO-DATE	AVAILABLE	...
deployment.apps/project-23-api	3/3	3	3	...

To note: we see another *Pod* here called `holy-api` which is part of another section. This is often the case in the provided scenarios, so be careful to only manipulate the resources you need to. Just like in the real world and in the exam.

Next we use `nginx:alpine` and `curl` to check if one *Pod* is accessible on port 80:

```
→ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 10.12.2.15
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload  Total   Spent    Left   Speed

<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
...
```

We could also use `busybox` and `wget` for this:

```
→ k run tmp --restart=Never --rm --image=busybox -i -- wget -O- 10.12.2.15
Connecting to 10.12.2.15 (10.12.2.15:80)
writing to stdout
-              100% |*****|          612  0:00:00 ETA
written to stdout
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
```

Now that we're sure the *Deployment* works we can continue with altering the provided yaml:

```
cp /opt/course/p1/project-23-api.yaml /opt/course/p1/project-23-api-new.yaml
vim /opt/course/p1/project-23-api-new.yaml
```

Add the liveness-probe to the yaml:

```
# /opt/course/p1/project-23-api-new.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: project-23-api
  namespace: pluto
spec:
  replicas: 3
  selector:
    matchLabels:
      app: project-23-api
  template:
    metadata:
      labels:
        app: project-23-api
    spec:
```

```

volumes:
- name: cache-volume1
  emptyDir: {}
- name: cache-volume2
  emptyDir: {}
- name: cache-volume3
  emptyDir: {}
containers:
- image: httpd:2.4-alpine
  name: httpd
  volumeMounts:
  - mountPath: /cache1
    name: cache-volume1
  - mountPath: /cache2
    name: cache-volume2
  - mountPath: /cache3
    name: cache-volume3
  env:
  - name: APP_ENV
    value: "prod"
  - name: APP_SECRET_N1
    value: "IO=a4L/XkRdvN8jM=Y+"
  - name: APP_SECRET_P1
    value: "-7PA0_Z]>{pwa43r) __"
  livenessProbe:                # add
    tcpSocket:                  # add
      port: 80                  # add
    initialDelaySeconds: 10      # add
    periodSeconds: 15           # add

```

Then let's apply the changes:

```
k -f /opt/course/p1/project-23-api-new.yaml apply
```

Next we wait 10 seconds and confirm the *Pods* are still running:

```
→ k -n pluto get pod
```

NAME	READY	STATUS	RESTARTS	AGE
holy-api	1/1	Running	0	144m
project-23-api-5b4579fd49-8knh8	1/1	Running	0	90s
project-23-api-5b4579fd49-cbgph	1/1	Running	0	88s
project-23-api-5b4579fd49-tcfq5	1/1	Running	0	86s

We can also check the configured liveness-probe settings on a *Pod* or the *Deployment*:

```
→ k -n pluto describe pod project-23-api-5b4579fd49-8knh8 | grep Liveness
Liveness:  tcp-socket :80 delay=10s timeout=1s period=15s #success=1 #failure=3

→ k -n pluto describe deploy project-23-api | grep Liveness
Liveness:  tcp-socket :80 delay=10s timeout=1s period=15s #success=1 #failure=3
```

Preview Question 2

Solve this question on instance: `ssh ckad9043`

Team Sun needs a new *Deployment* named `sunny` with 4 replicas of image `nginx:1.17.3-alpine` in *Namespace* `sun`. The *Deployment* and its *Pods* should use the existing *ServiceAccount* `sa-sun-deploy`.

Expose the *Deployment* internally using a ClusterIP *Service* named `sun-srv` on port 9999. The nginx containers should run as default on port 80. The management of Team Sun would like to execute a command to check that all *Pods* are running on occasion. Write that command into file `/opt/course/p2/sunny_status_command.sh`. The command should use `kubect1`.

Answer

```
k -n sun create deployment -h #help

k -n sun create deployment sunny --image=nginx:1.17.3-alpine --dry-run=client -oyaml > p2_sunny.yaml

vim p2_sunny.yaml
```

Then alter its yaml to include the requirements:

```
# p2_sunny.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: sunny
  name: sunny
  namespace: sun
spec:
  replicas: 4 # change
  selector:
    matchLabels:
      app: sunny
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: sunny
    spec:
      serviceAccountName: sa-sun-deploy # add
      containers:
      - image: nginx:1.17.3-alpine
        name: nginx
        resources: {}
  status: {}
```

Now create the yaml and confirm it's running:

```
→ k create -f p2_sunny.yaml
deployment.apps/sunny created
```

```
→ k -n sun get pod
```

NAME	READY	STATUS	RESTARTS	AGE
0509649a	1/1	Running	0	149m
0509649b	1/1	Running	0	149m
1428721e	1/1	Running	0	149m
...				
sunny-64df8dbdbb-9mxbw	1/1	Running	0	10s
sunny-64df8dbdbb-mp5cf	1/1	Running	0	10s
sunny-64df8dbdbb-pggdf	1/1	Running	0	6s
sunny-64df8dbdbb-zvqth	1/1	Running	0	7s

Confirmed, the AGE column is always in important information about if changes were applied. Next we expose the *Pods* by created the *Service*:

```
k -n sun expose -h # help
k -n sun expose deployment sunny --name sun-srv --port 9999 --target-port 80
```

Using expose instead of `kubectl create service clusterip` is faster because it already sets the correct selector-labels. The previous command would produce this yaml:

```
# k -n sun expose deployment sunny --name sun-srv --port 9999 --target-port 80
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: sunny
  name: sun-srv          # required by task
spec:
  ports:
    - port: 9999        # service port
      protocol: TCP
      targetPort: 80     # target port
  selector:
    app: sunny          # selector is important
status:
  loadBalancer: {}
```

Let's test the *Service* using `wget` from a temporary *Pod*:

```
→ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 sun-srv.sun:9999
Connecting to sun-srv.sun:9999 (10.23.253.120:9999)
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
...
```

Because the *Service* is in a different *Namespace* as our temporary *Pod*, it is reachable using the names `sun-srv.sun` or fully: `sun-srv.sun.svc.cluster.local`.

Finally we need a command which can be executed to check if all *Pods* are runing, this can be done with:

```
vim /opt/course/p2/sunny_status_command.sh
```

```
# /opt/course/p2/sunny_status_command.sh
kubect1 -n sun get deployment sunny
```

To run the command:

```
→ sh /opt/course/p2/sunny_status_command.sh
NAME      READY    UP-TO-DATE    AVAILABLE    AGE
sunny     4/4      4             4            13m
```

Preview Question 3

Solve this question on instance: `ssh ckad5601`

Management of EarthAG recorded that one of their *Services* stopped working. Dirk, the administrator, left already for the long weekend. All the information they could give you is that it was located in *Namespace* `earth` and that it stopped working after the latest rollout. All *Services* of EarthAG should be reachable from inside the cluster.

Find the *Service*, fix any issues and confirm it's working again. Write the reason of the error into file `/opt/course/p3/ticket-654.txt` so Dirk knows what the issue was.

Answer

First we get an overview of the resources in *Namespace* `earth`:

```
→ k -n earth get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/earth-2x3-api-584df69757-ngnwp	1/1	Running	0	116m
pod/earth-2x3-api-584df69757-ps8cs	1/1	Running	0	116m
pod/earth-2x3-api-584df69757-ww9q8	1/1	Running	0	116m
pod/earth-2x3-web-85c5b7986c-48vjt	1/1	Running	0	116m
pod/earth-2x3-web-85c5b7986c-6mqmb	1/1	Running	0	116m
pod/earth-2x3-web-85c5b7986c-6vj1l	1/1	Running	0	116m
pod/earth-2x3-web-85c5b7986c-fnkbp	1/1	Running	0	116m
pod/earth-2x3-web-85c5b7986c-pjm5m	1/1	Running	0	116m
pod/earth-2x3-web-85c5b7986c-pwfvj	1/1	Running	0	116m
pod/earth-3cc-runner-6cb6cc6974-8wm5x	1/1	Running	0	116m
pod/earth-3cc-runner-6cb6cc6974-9fx8b	1/1	Running	0	116m
pod/earth-3cc-runner-6cb6cc6974-b9nrv	1/1	Running	0	116m
pod/earth-3cc-runner-heavy-6bf876f46d-b47vq	1/1	Running	0	116m
pod/earth-3cc-runner-heavy-6bf876f46d-mrzqd	1/1	Running	0	116m
pod/earth-3cc-runner-heavy-6bf876f46d-qkd74	1/1	Running	0	116m
pod/earth-3cc-web-6bfdf8b848-f74cj	0/1	Running	0	116m
pod/earth-3cc-web-6bfdf8b848-n4z7z	0/1	Running	0	116m
pod/earth-3cc-web-6bfdf8b848-rcmxs	0/1	Running	0	116m
pod/earth-3cc-web-6bfdf8b848-x1467	0/1	Running	0	116m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/earth-2x3-api-svc	ClusterIP	10.3.241.242	<none>	4546/TCP	116m
service/earth-2x3-web-svc	ClusterIP	10.3.250.247	<none>	4545/TCP	116m

service/earth-3cc-web	ClusterIP	10.3.243.24	<none>	6363/TCP	116m
-----------------------	-----------	-------------	--------	----------	------

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/earth-2x3-api	3/3	3	3	116m
deployment.apps/earth-2x3-web	6/6	6	6	116m
deployment.apps/earth-3cc-runner	3/3	3	3	116m
deployment.apps/earth-3cc-runner-heavy	3/3	3	3	116m
deployment.apps/earth-3cc-web	0/4	4	0	116m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/earth-2x3-api-584df69757	3	3	3	116m
replicaset.apps/earth-2x3-web-85c5b7986c	6	6	6	116m
replicaset.apps/earth-3cc-runner-6cb6cc6974	3	3	3	116m
replicaset.apps/earth-3cc-runner-heavy-6bf876f46d	3	3	3	116m
replicaset.apps/earth-3cc-web-6895587dc7	0	0	0	116m
replicaset.apps/earth-3cc-web-6bfdf8b848	4	4	0	116m
replicaset.apps/earth-3cc-web-d49645966	0	0	0	116m

First impression could be that all *Pods* are in status `RUNNING`. But looking closely we see that some of the *Pods* are not ready, which also confirms what we see about one *Deployment* and one *ReplicaSet*. This could be our error to further investigate.

Another approach could be to check the *Services* for missing endpoints:

```
→ k -n earth get ep
```

NAME	ENDPOINTS	AGE
earth-2x3-api-svc	10.0.0.10:80,10.0.1.5:80,10.0.2.4:80	116m
earth-2x3-web-svc	10.0.0.11:80,10.0.0.12:80,10.0.1.6:80 + 3 more...	116m
earth-3cc-web		

Service `earth-3cc-web` doesn't have endpoints. This could be a selector/label misconfiguration or the endpoints are actually not available/ready.

Checking all *Services* for connectivity should show the same (this step is optional and just for demonstration):

```
→ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 earth-2x3-api-svc.earth:4546
...
<html><body><h1>It works!</h1></body></html>
```

```
→ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 earth-2x3-web-svc.earth:4545
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload   Total   Spent    Left  Speed
100    45    100    45     0     0   5000      0  --:--:--  --:--:--  --:--:--   5000
<html><body><h1>It works!</h1></body></html>
```

```
→ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 earth-3cc-web.earth:6363
If you don't see a command prompt, try pressing enter.
  0     0     0     0     0     0     0     0  --:--:--  0:00:05  --:--:--     0
curl: (28) Connection timed out after 5000 milliseconds
pod "tmp" deleted
pod default/tmp terminated (Error)
```

Notice that we use here for example `earth-2x3-api-svc.earth`. We could also spin up a temporary *Pod* in *Namespace* `earth` and connect directly to `earth-2x3-api-svc`.

We get no connection to `earth-3cc-web.earth:6363`. Let's look at the *Deployment* `earth-3cc-web`. Here we see that the requested amount of replicas is not available/ready:


```
→ k -n earth get deploy earth-3cc-web
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
earth-3cc-web	0/4	4	0	7m18s

To continue we check the *Deployment* yaml for some misconfiguration:

```
k -n earth edit deploy earth-3cc-web
```

```
# k -n earth edit deploy earth-3cc-web
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
...
  generation: 3                                # there have been rollouts
  name: earth-3cc-web
  namespace: earth
...
spec:
...
  template:
    metadata:
      creationTimestamp: null
      labels:
        id: earth-3cc-web
    spec:
      containers:
      - image: nginx:1.16.1-alpine
        imagePullPolicy: IfNotPresent
        name: nginx
        readinessProbe:
          failureThreshold: 3
          initialDelaySeconds: 10
          periodSeconds: 20
          successThreshold: 1
          tcpSocket:
            port: 82                                # this port doesn't seem to be right, should be 80
          timeoutSeconds: 1
...

```

We change the readiness-probe port, save and check the *Pods*:

```
→ k -n earth get pod -l id=earth-3cc-web
```

NAME	READY	STATUS	RESTARTS	AGE
earth-3cc-web-d49645966-52vb9	0/1	Running	0	6s
earth-3cc-web-d49645966-5tts6	0/1	Running	0	6s
earth-3cc-web-d49645966-db5gp	0/1	Running	0	6s
earth-3cc-web-d49645966-mk7gr	0/1	Running	0	6s

Running, but still not in ready state. Wait 10 seconds (initialDelaySeconds of readinessProbe) and check again:

```
→ k -n earth get pod -l id=earth-3cc-web
```

NAME	READY	STATUS	RESTARTS	AGE
earth-3cc-web-d49645966-52vb9	1/1	Running	0	32s
earth-3cc-web-d49645966-5tts6	1/1	Running	0	32s
earth-3cc-web-d49645966-db5gp	1/1	Running	0	32s
earth-3cc-web-d49645966-mk7gr	1/1	Running	0	32s

Let's check the service again:

```
→ k run tmp --restart=Never --rm -i --image=nginx:alpine -- curl -m 5 earth-3cc-web.earth:6363
% Total      % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left   Speed

100    612    100    612      0      0  55636      0 --:--:-- --:--:-- --:--:--  55636

<!DOCTYPE html>

<html>

<head>

<title>Welcome to nginx!</title>

<style>

    body {

        width: 35em;

        margin: 0 auto;

        font-family: Tahoma, Verdana, Arial, sans-serif;

    }

</style>

</head>

<body>

<h1>Welcome to nginx!</h1>

...

```

We did it! Finally we write the reason into the requested location:

```
vim /opt/course/p3/ticket-654.txt
```

```
# /opt/course/p3/ticket-654.txt
yo Dirk, wrong port for readinessProbe defined!
```

CKAD Tips Kubernetes 1.33

In this section we'll provide some tips on how to handle the CKAD exam and browser terminal.

Knowledge

- Study all topics as proposed in the curriculum until you feel comfortable with all
- Learn and Study the in-browser scenarios on <https://killercoda.com/killer-shell-ckad>
- Read this and do all examples: <https://kubernetes.io/docs/concepts/cluster-administration/logging>
- Understand Rolling Update Deployment including maxSurge and maxUnavailable
- Do 1 or 2 test sessions with this CKAD Simulator. Understand the solutions and maybe try out other ways to achieve the same
- Be fast and breathe `kubect1`

CKAD Preparation

Read the Curriculum

<https://github.com/cncf/curriculum>

Read the Handbook

<https://docs.linuxfoundation.org/tc-docs/certification/lf-handbook2>

Read the important tips


Read the FAQ

<https://docs.linuxfoundation.org/tc-docs/certification/faq-cka-ckad>

Kubernetes documentation

Get familiar with the Kubernetes documentation and be able to use the search. Allowed resources are:

- <https://kubernetes.io/docs>
- <https://kubernetes.io/blog>
- <https://helm.sh/docs>

 Verify the list [here](#)

The Exam UI / Remote Desktop

The real exam, as well as the simulator, provides a Remote Desktop (XFCE) on Ubuntu/Debian. Coming from OSX/Windows there will be changes in copy&paste for example.

Official Information

[ExamUI: Performance Based Exams](#)


Lagging

There could be some lagging, definitely make sure you are using a good internet connection because your webcam and screen are transferring all the time.

Kubectl autocomplete and commands

The following are installed or pre-configured, verify the list [here](#):

- `kubectl` with `k` alias and Bash autocomplete
- `yq` or YAML processing
- `curl` and `wget` for testing web services
- `man` and man pages for further documentation

 You're allowed to install tools, like `tmux` for terminal multiplexing or `jq` for JSON processing

Copy & Paste

Copy and pasting will work like normal in a Linux Environment:

What always works: copy+paste using right mouse context menu What works in Terminal: Ctrl+Shift+c and Ctrl+Shift+v What works in other apps like Firefox: Ctrl+c and Ctrl+v

Score

There are 15-20 questions in the exam. Your results will be automatically checked according to the handbook. If you don't agree with the results you can request a review by contacting the Linux Foundation Support.

Notepad & Flagging Questions

You can flag questions to return to later. This is just a marker for yourself and won't affect scoring. You also have access to a simple notepad in the browser which can be used to store any kind of plain text. It might make sense to use this and write down

additional information about flagged questions. Instead of using the notepad you could also open Mousepad (XFCE application inside the Remote Desktop) or create a file with Vim.

VSCodium

You can use VSCodium to edit files and you can also use its terminal to run commands. You're not allowed to install any VSCodium extensions.

Servers

Each question needs to be solved on a specific instance other than your main terminal. You'll need to connect to the correct instance via ssh, the command is provided before each question.

PSI Bridge

Starting with [PSI Bridge](#):

- The exam will now be taken using the PSI Secure Browser, which can be downloaded using the newest versions of Microsoft Edge, Safari, Chrome, or Firefox
- Multiple monitors will no longer be permitted
- Use of personal bookmarks will no longer be permitted

The new ExamUI includes improved features such as:

- A remote desktop configured with the tools and software needed to complete the tasks
- A timer that displays the actual time remaining (in minutes) and provides an alert with 30, 15, or 5 minute remaining
- The content panel remains the same (presented on the Left Hand Side of the ExamUI)

Read more [here](#).

Terminal Handling

Bash Aliases

In the real exam, each question has to be solved on a different instance to which you connect via ssh. This means it's not advised to configure bash aliases because they wouldn't be available on the instances accessed by ssh.

Be fast

Use the `history` command to reuse already entered commands or use even faster history search through **Ctrl +r**.

If a command takes some time to execute, like sometimes `kubect1 delete pod x`. You can put a task in the background using **Ctrl +z** and pull it back into foreground running command `fg`.

You can delete *Pods* fast with:

```
k delete pod x --grace-period 0 --force
```

Vim

Be great with vim.

Settings

In case you face a situation where vim is not configured properly and you face for example issues with pasting copied content you should be able to configure via `~/.vimrc` or by entering manually in vim settings mode:

```
set tabstop=2
set expandtab
set shiftwidth=2
```

The `expandtab` option makes sure to use spaces for tabs.

Note that changes in `~/.vimrc` will not be transferred when connecting to other instances via ssh.

Toggle vim line numbers

When in `vim` you can press **Esc** and type `:set number` or `:set nonumber` followed by **Enter** to toggle line numbers. This can be useful when finding syntax errors based on line - but can be bad when wanting to mark© by mouse. You can also just jump to a line number with **Esc** `:22` + **Enter**.

Copy&Paste

Get used to copy/paste/cut with vim:

```
Mark lines: Esc+V (then arrow keys)
Copy marked lines: y
Cut marked lines: d
Paste lines: p or P
```

Indent multiple lines

To indent multiple lines press **Esc** and type `:set shiftwidth=2`. First mark multiple lines using `Shift v` and the up/down keys. Then to indent the marked lines press `>` or `<`. You can then press `.` to repeat the action.

[About](#)

[FAQ](#)

[Support](#)

[Store](#)

[Pricing](#)

[Legal Notice / Impressum](#)

[Privacy Policy / Datenschutz](#)

[Terms / AGB](#)

CONTENT

[CKS](#)

[CKA](#)

[CKAD](#)

[LFCS](#)

LINKS

[Killercoda](#)

[Kim Wuestkamp](#)

